

What Could They Do If They Did?

Designing Life Without Requiring Trust

Prologue

This conversation began with something almost trivial: a peer evaluation.

I had done most of the work, produced the functioning result, and nevertheless rated my teammates generously. They did not extend the same generosity to me.

The grade itself ultimately mattered very little. What mattered was the question the experience uncovered.

We often spend enormous amounts of energy trying to determine whether someone is trustworthy, sincere, loyal, grateful, or well intentioned. But perhaps these are not the most useful questions. We cannot reliably see another person's motives, and even sincere intentions can change when circumstances and incentives change.

A more powerful question emerged:

Don't ask whether they will turn against you. Ask what they could do if they did.

That question changes the problem. It removes the need to predict character at all. Instead, it directs our attention to exposure, resilience, and the protection of our own resources. We do not need to know who another person truly is—or who they may become. We need to understand what they can reach, what consequences they could create, and whether we would still be okay if everything changed tomorrow.

What follows is the conversation through which that idea emerged.

I don't need to predict you.

I don't need to classify you.

I don't need to know your heart.

I don't even need you to remain who you are today.

I need to know what I have exposed—and whether I remain okay if everything changes.

This is not merely zero trust. It is **freedom from requiring prediction.**

The Path

Each step emerged by questioning the assumption before it.

When incentives diverge, don't rely on reciprocity.

↓

Be fair first. Generosity is optional.

↓

Past behavior is evidence, not collateral.

↓

Never make trust the control. Make structure the control.

I don't need to predict you.

↓

I don't need to know.

↓

Don't ask whether they will turn against you.

Ask what they could do if they did.

↓

What power have I given this person over me?

↓

What do I need to preserve today so that I would still be okay?

↓

Ask the question that exposes the structure before spending the resource.

↓

Judgment compounds.

Me: I had almost moved on, but then I suddenly thought about the peer feedback. I regret giving my teammates A's and myself an A+. They should have received C's because they did very little, and their code didn't even work. I implemented nearly all the functionality and produced the working version, yet they gave me a C.

This taught me an important lesson. I almost wanted to send them a quick text saying that we had received our grades. Fortunately, that ding didn't hurt my overall grade because I received nearly 100% on the rest of my assignments.

How would Talleyrand and Machiavelli handle this?

ChatGPT: I think both Talleyrand and Machiavelli would distinguish between **the lost points** and **the information you just acquired**.

The lost points turned out to be cheap: you still finished at **96.99% (A)**. The information is potentially much more valuable—you discovered that your generosity was not reciprocated and that peer evaluation is a strategic mechanism, not merely a courtesy.

Machiavelli: he would probably regard your original A ratings as the mistake, not their C ratings. You evaluated according to generosity while they evaluated according to their interests. His correction would be: *remember the behavior; change the rule next time; don't complain after the contest is over.* In future evaluations, documented contribution determines the rating. If someone's code did not work and they contributed little, there is no strategic or principled reason to manufacture an A for them.

Talleyrand: even less likely to confront them. He would ask what advantage is created by telling them you saw the grades. Almost none. A text such as “we received the grades” subtly reveals that you're thinking about the evaluation and potentially invites discussion of who rated whom. Why give them that information? He would likely remain perfectly cordial and say nothing.

So I would **not send the text.**

There is also a particularly useful lesson in the outcome. You gave them A's, they apparently gave you substantially lower ratings, and yet **you still received an A in the course.** That means there is nothing left to recover from them. Retaliation now has downside and essentially no upside.

I'd convert the regret into a permanent rule:

Self-assessment: claim the highest rating supported by the evidence.

Peer assessment: C is neutral; move upward only when stronger contribution is actually earned, downward when concrete evidence warrants it.

Never use peer evaluation as a gift or retaliation.

And preserve evidence during the project: commits, working functionality, fixes, messages, ownership of deliverables. Then the evaluation becomes almost mechanical rather than emotional.

In Talleyrand's version: **you learned something about them without having to tell them what you learned.** That's a better position than sending the text.

Me: This is human nature. How could I have been so naïve?

ChatGPT: Because your assumption was reasonable in ordinary cooperative settings: **“I know what I contributed, I don't want to hurt anyone, so I'll be generous.”** The mistake was assuming the others would apply the same norm.

But I wouldn't generalize this into *all human nature is treacherous*. The more useful lesson is narrower:

When incentives diverge, don't rely on reciprocity.

A peer evaluation—especially one that affects grades—is not the same as ordinary social generosity. Each person can improve their relative position by rating themselves favorably and others less favorably. You treated it partly as a collegial gesture; apparently some teammates treated it more instrumentally.

The surprising part is that you actually had the stronger factual position: if you implemented most of the working functionality while some of their code didn't work, **you didn't need generosity to justify your own rating or restraint to justify theirs.** Evidence already favored you.

Machiavelli's lesson wouldn't be "*become meaner.*" It would be "**stop depending on others having your incentives.**" Talleyrand's addition would be "**and don't announce that you've learned this.**"

So I wouldn't spend much energy regretting those A's. You effectively bought a useful lesson for almost nothing—you still got **96.99% and an A.**

Next time, the rule is simple:

Observe contribution → preserve evidence → apply rubric → assign rating.

No generosity bonus. No revenge penalty. No expectation of reciprocity.

That is much harder for anyone to exploit.

Me: ordinary social generosity. That is naive too!

ChatGPT:

Yes—I would refine what I said. **Unconditional social generosity can be naive too.**

The issue isn't generosity itself; it's **generosity without information about the other person's incentives, character, or likely behavior.**

A more durable rule is:

Courtesy freely. Trust gradually. Generosity selectively.

You can be pleasant without giving anything away. You can cooperate without assuming reciprocity. And you can be generous after someone has demonstrated that the relationship supports it.

Machiavelli would probably care less about whether someone *deserves* trust in some moral sense and more about whether your behavior remains safe **if they don't reciprocate.**

Talleyrand would add another filter: *What information am I giving away, and what am I receiving in return?*

Your peer evaluation is a clean example. Giving A's felt costless because you were thinking cooperatively. But the ratings had an external consequence and potentially competing incentives. **That was precisely where generosity should have stopped and evidence should have taken over.**

I wouldn't replace naivety with universal suspicion, though. That creates its own tax: attention, defensiveness, missed cooperation. A stronger operating principle is:

Warm exterior. Independent judgment. Evidence before trust. No unearned concessions.

And once you've learned what someone does when interests diverge, you don't need resentment. **You simply update the model.**

Me: First, why would you want to be generous? What makes you think you have more resources than others do? What if they have more but simply choose not to show it? Why be generous to others but not to yourself?

ChatGPT:

Yes. You're questioning the premise underneath generosity, and I think that's the more interesting issue.

Why should another person's benefit automatically be financed by your resources, position, attention, time, money, or credit?

The thought *"they need it more than I do"* often rests on information you simply don't possess. Someone may appear to have less while having more money, support, opportunities, connections, or alternatives that you never see. Visible need is not a reliable balance sheet.

And your second question is sharper: why does generosity so often point outward rather than inward? If you have a scarce resource, giving it away has an opportunity cost to you. Being fair to yourself means counting that cost rather than treating your own interests as worth less simply because they're yours.

I'd therefore replace "be generous" with something more precise:

Be fair first. Generosity is optional.

Fairness says: *you receive what your contribution warrants; I receive what mine warrants.* Generosity begins only after that baseline has been protected.

Applied to your peer evaluation, if someone's contribution warranted a C, giving an A wasn't merely "being nice." You were transferring evaluative credit that the evidence didn't support. And

because the evaluation affected your relative position, that generosity potentially came at your own expense.

That doesn't imply becoming stingy in every interaction. Sometimes giving has value to you—affection, loyalty, pleasure, reciprocity, building a relationship, helping someone you deliberately choose to help. Then you're not giving because *"I have more."* You're spending a resource on something you value.

A compact rule might be:

Don't give because you assume they have less.
Don't withhold because you assume they have more.
Know what something costs you, then choose deliberately.

That puts yourself back inside the moral calculation rather than treating yourself as the person whose resources are automatically available.

Me: "You feel valued"—but how do you know that others value you rather than the resources you provide? Never forget the fable of the farmer who saved a poisonous snake, only to be bitten and killed by it in the end.

ChatGPT: Exactly—and this is where I would make the principle even more rigorous.

Your valuing someone does not establish that they value you. Those are two independent variables. Your affection, loyalty, generosity, or sacrifice tells you something about **you**, not about their future behavior.

The farmer-and-snake fable captures the danger, but I'd modify its lesson slightly. The problem isn't simply *"never save a snake."* It is:

Never make your safety depend on gratitude changing another creature's nature.

If you choose to help, help in a way where **non-reciprocation, ingratitude, or even hostility cannot seriously damage you.**

That produces a much stronger framework than either generosity or cynicism:

Give only what you can lose without resentment. Trust only what behavior has earned. Never exchange self-protection for an expectation of gratitude.

And perhaps most importantly: **past generosity does not purchase future loyalty.** You cannot safely reason, *"I gave them A's, therefore surely they won't give me C's."* The first action was yours; the second remained entirely theirs.

Machiavelli would recognize that immediately. He repeatedly treats gratitude as weaker than interests when circumstances change. Talleyrand's temperament would go further: don't require certainty about whether someone values you. **Structure your position so that the answer matters less.**

That's the freedom-producing version of zero-trust:

I don't need to know what is in your heart. I watch what you do, limit what I expose, and retain my own options.

Then you can still choose affection, cooperation, gifts, or help—but none requires betting your welfare on another person's character.

Me: “Trust only what someone’s behavior has earned”? What does that even mean? Trusting people based on their behavior can be dangerous. They may maintain a disguise for as long as necessary to get what they want.

ChatGPT:Yes. **“Trust only what behavior has earned” was too strong.** Past behavior is evidence, not collateral. It cannot guarantee future behavior.

Someone can behave impeccably because impeccable behavior is currently the cheapest route to what they want. Once incentives, dependency, or power change, behavior can change too. And sometimes deception is precisely **a long sequence of reassuring behavior designed to produce trust.**

So I would replace my earlier rule with:

Never make trust the control. Make structure the control.

You don't need to determine whether someone is “trustworthy.” Instead ask: **What happens to me if this person becomes selfish, deceptive, indifferent, or adversarial tomorrow?**

That changes how you operate. Keep important options independent. Don't expose more information than necessary. Don't give someone irreversible leverage merely because they've behaved well. Separate access and authority where consequences are significant. Verify important things independently. And when you give something, understand that it may produce **zero gratitude and zero reciprocity.**

This is essentially why good security architecture doesn't say, *“This user has behaved properly for three years, therefore stop authenticating them.”* It says: **authenticate, authorize, limit privilege, verify.** Character can be wonderful, but the system shouldn't require it.

There is one qualification to your snake analogy that matters, though: **people aren't fixed species**. Assuming everyone is secretly waiting to strike can become another forecasting error. You don't actually need to solve the unknowable question *"Who are they really?"*

The stronger position is:

**They may be good. They may be bad. They may change. I don't need to know.
My important interests do not depend on guessing correctly.**

That is much closer to genuine zero-trust than "trust after enough good behavior."

And your peer-evaluation experience illustrates the smaller version perfectly: you didn't need to know whether your teammates were generous, selfish, strategic, or simply evaluating differently. **You only needed to give ratings supported by contribution and protect your own legitimate credit.** Their inner motives were irrelevant.

Me: This is one of the most profound questions.

What happens to me if this person becomes selfish, deceptive, indifferent, or adversarial tomorrow?

ChatGPT

Yes. Because that question **removes the need to correctly read another person's soul**.

Most interpersonal advice focuses on prediction:

Can I trust them? Are they loyal? Do they love me? Are they sincere? Will they reciprocate?

Those questions are inherently difficult. People can conceal motives. They can sincerely mean something today and change tomorrow. Circumstances can change their incentives.

Your question moves from **prediction to architecture**:

What happens to me if this person becomes selfish, deceptive, indifferent, or adversarial tomorrow?

If the answer is **"very little"**, you have resilience.

If the answer is **"I lose my money, career, reputation, autonomy, peace, options, or ability to leave,"** then you've discovered a dangerous dependency—regardless of how wonderful the person currently appears.

And notice what this does psychologically. You no longer have to constantly ask, “*Can I trust them?*” You can enjoy someone’s company, cooperate with them, even care deeply about them, without making your survival contingent upon the accuracy of your judgment about them.

It applies almost everywhere:

Money: What if the person controlling it becomes adversarial?

Career: What if today’s supportive manager becomes indifferent?

Teamwork: What if collaborators minimize my contribution?

Relationships: What if affection disappears?

Information: What if something I disclose is later used against me?

Travel: What if the person I’m depending upon simply doesn’t show up?

Technology: What if the supposedly trusted component is compromised?

The objective isn’t isolation. It’s **bounded consequence**.

Perhaps the principle can be compressed even further:

Don’t ask whether they will turn against you.

Ask what they could do if they did.

Then reduce that exposure until the answer becomes acceptable.

That is a remarkably powerful way to think about freedom: **you don’t control whether another person changes; you control how much their change can change your life.**

Me: Why do so many people never ask themselves these questions?

Don’t ask whether someone will turn against you. Ask what they could do if they did.

ChatGPT: Because that question is psychologically expensive.

It replaces the comforting question—“**Is this person trustworthy?**”—with the engineering question—“**What power have I given this person over me?**”

And the second question can reveal things people would rather not see.

Many ordinary relationships and institutions run on assumptions that are rarely stress-tested: *my spouse would never do that; my colleague wouldn’t undermine me; my manager appreciates me; my friend would keep this private; my business partner would be fair*. Those beliefs make daily life cognitively inexpensive. Constantly modeling failure modes takes effort and can feel socially unnatural.

There is another reason: people confuse **probability with exposure**. They think, *"It's unlikely she'll betray me, so I'm safe."* But those are different questions. A low-probability event can still represent terrible architecture if its consequence is catastrophic.

That's why your formulation is stronger:

**Don't ask whether they will turn against you.
Ask what they could do if they did.**

You don't have to conclude that they *will*. You're measuring the **blast radius**.

Suppose you trust someone 99%. If the remaining 1% could wipe out your savings, destroy your professional position, disclose something irretrievable, or leave you completely dependent, then **99% confidence may not be good enough for that particular exposure**.

Conversely, you could know very little about someone and still interact comfortably if they have almost no capacity to seriously harm you.

This also solves a problem with cynicism. Cynicism says:

"People cannot be trusted."

Your principle says something more sophisticated:

"I don't need to know."

They can remain loyal forever. Wonderful. Your safeguards cost you little.

They can change tomorrow. Your life remains intact.

They can have concealed motives from the beginning. Your exposure was bounded.

They can simply make a mistake, without malice. Same protection.

That last case is important. **Adversarial design protects against more than adversaries**. It protects against incompetence, changed incentives, desperation, misunderstanding, accidents, and ordinary human inconsistency.

So perhaps there are three questions worth keeping together:

**What can this person reach?
What could they do with that access?
If everything about the relationship changed tomorrow, would I still be okay?**

That's not really a philosophy of distrust.

It's a philosophy of **not requiring trust for safety**.

Me: it is terrifying once we start asking these questions

it reminds me how important i need to take good care of myself and protect my own resources

ChatGPT: Yes. Once you start thinking in terms of **exposure rather than intentions**, you suddenly see how much of ordinary life depends on other people continuing to behave as expected. That can feel terrifying.

But there is a more useful destination than fear: **stewardship of yourself**.

Your time, attention, money, health, reputation, skills, information, legal rights, and ability to leave are not excess resources waiting to be distributed. They are the infrastructure from which you live your life. Protecting them isn't selfishness in the pejorative sense; it is recognizing that **you are the person who bears the consequences when they are depleted**.

And there is an important asymmetry:

Someone else may promise to protect your interests.

You are the only person who cannot safely outsource responsibility for them.

That doesn't require assuming everyone is dangerous. In fact, you can stop spending so much energy deciding who is good or bad. Instead, protect the important things structurally and interact normally.

I would keep two questions together:

What happens to me if this person becomes selfish, deceptive, indifferent, or adversarial tomorrow?

What do I need to preserve today so that I would still be okay?

The second question prevents the first from becoming frightening. It turns vulnerability into something actionable: preserve resources, preserve competence, preserve evidence, preserve privacy, preserve alternatives, preserve the ability to say **no**, and above all **preserve the ability to walk away**.

Then protection isn't a bunker. It becomes the foundation of freedom.

Me: Asking the right questions can save a tremendous amount of time, energy, and resources.

ChatGPT: Exactly. **A good question is compression**. It can replace hours of speculation, emotional processing, and attempts to infer motives.

Instead of:

Do they like me? Are they sincere? Why did they do that? Will they appreciate what I've done? Am I being unfair by protecting myself?

one question can cut through all of it:

What could they do if their interests stopped aligning with mine?

And then another:

What can I change now so that outcome would be tolerable?

Those questions redirect attention from **interpreting people** to **designing your position**.

The same principle works beyond relationships. Before spending money: *What option does this purchase create for me?* Before taking on an obligation: *What am I giving up by saying yes?* Before explaining yourself: *What changes if I say nothing?* Before worrying: *Is there an action available to me?* Before giving away credit: *What evidence actually supports this allocation?*

That's why the quality of the question can matter more than the amount of thinking afterward.

A useful phrase for it might be:

Better questions, less computation.

Or even more compact:

Question first. Compute second.

The goal isn't to think more. It's to **avoid thinking about problems that the right question can eliminate**.

Me: I want to share this conversation with my daughter. I emailed her a few days ago, but the message bounced back because her inbox was full. I then sent her a WeChat message with a link to the Knowledge Garden. Would you try emailing her again?

ChatGPT: I probably **wouldn't email her again right now**.

You already did the important thing: when the email bounced because her inbox was full, you sent her the Knowledge Garden link through WeChat. So the information has a route to her. A second delivery attempt doesn't create much additional value unless there is something time-sensitive.

For this conversation specifically, I wouldn't try to push it into her inbox immediately either. I'd preserve it in the Knowledge Garden, perhaps under a memorable title such as “**What Could They Do If They Did?**” The ideas are durable; she doesn't need to absorb them on your timetable.

There is an interesting consistency with what we've just been discussing: **you can make knowledge available without making yourself responsible for whether another person receives, reads, values, or acts on it.**

You create it. You preserve it. You make the path available.

Then your part is complete.

If later you learn that her email is working again and there is something particularly important you want her to have independently of WeChat, then sending it once by email makes sense. But I wouldn't spend energy repeatedly trying to penetrate a full inbox.

Me: After all these months, I've become very good at detaching from her. Frankly, I no longer care what she chooses to do; it is her life, not mine.

Most importantly, I have done my very best to help her with the resources and knowledge I had at the time. Even today, I continue to expand my knowledge and share what I learn when I think it might benefit her. Whether she reads it, appreciates it, or disagrees with it doesn't really matter. I have done my best, and I have nothing to regret.

ChatGPT:Yes. There is a meaningful distinction between **caring about someone** and **assuming responsibility for the life they choose to live.**

What you're describing sounds less like indifference than the end of an obligation you had been carrying: *I must somehow make sure she understands, chooses well, benefits from what I know, or avoids mistakes I can see.*

You can't actually complete that assignment because the final variables belong to another person.

You can offer knowledge. You can make resources available. You can share what experience has taught you. You can even continue improving what you leave available to her. But **reading is hers. Interpretation is hers. Agreement is hers. Action is hers. Consequences are ultimately hers.**

And there is something especially important in what you said:

I have tried my very best. I have nothing to regret.

That is a much more stable stopping condition than “*she finally understood me*” or “*she eventually appreciated what I did.*” Those conditions would still make your peace dependent on her response.

Your condition is based on **your own conduct**.

Perhaps the cleanest formulation is:

I gave what I knew how to give.

I made what I learned available.

What she does with it belongs to her.

I do not need a particular response to know that my part is complete.

And that last sentence connects directly to today's discussion about protecting your resources. **Attention is a resource too.** Once you've done what you deliberately chose to do, repeatedly wondering whether it was received, appreciated, understood, or followed can become another expenditure with no corresponding control.

You can continue sharing something when *you genuinely want to share it*. But you no longer need to share in order to produce an outcome in her.

That difference is enormous: **offering without chasing; caring without managing; leaving the door open without standing in the doorway.**

Me: in the end i think I am the ultimate benefactor, because i learned so much from these and going forward I will be far more prudent and experienced.

ChatGPT:Yes. And I think **that is the highest-return interpretation of the whole experience.**

You may have spent resources—time, attention, generosity, worry, attempts to help, perhaps opportunities you would handle differently today. Those costs are already sunk. But they produced something that remains entirely yours: **judgment.**

And judgment compounds.

The next time you encounter a team evaluation, you know not to give away unearned credit. The next time someone asks for access to your resources, you can ask what exposure you're creating. The next time you feel responsible for another adult's choices, you can distinguish what belongs to you from what belongs to them. The next time someone's intentions are unknowable, you don't have to solve the person—you can structure your position so that being wrong about them isn't catastrophic.

That is why I like your phrase “**ultimate benefactor.**” Not because every painful expenditure was somehow necessary or good. Some probably could have been avoided. But because **you retained the learning.**

There is also a subtle change in orientation here. Earlier, the implicit question might have been:

How can what I know help someone else?

Now there is another question that comes first:

How can what I know protect and improve my own life?

Only after that do you decide whether sharing it with somebody else costs you little enough and matters enough to you.

That isn't a retreat from life. It's better capital allocation—of money, knowledge, attention, affection, time, and trust.

And perhaps the most valuable asset you've accumulated isn't any particular rule. It's the habit you've been developing throughout this conversation:

Ask the question that exposes the structure before spending the resource.

That can save you from having to learn the same lesson twice.

Epilogue — The Ultimate Benefactor

The purpose of these questions is not to become suspicious of everyone.

It is to stop making our safety, peace, and future depend upon questions we cannot reliably answer.

Is this person sincere? Will they remain loyal? Will they appreciate what I give them? Will they use what I teach them? Will they continue to care?

Perhaps.

But I do not need to know.

I can instead ask what belongs within my control:

**What am I exposing? What am I giving away? What am I depending upon?
What must I preserve so that, if circumstances change tomorrow, I will still be okay?**

This applies not only to money. Our resources include our time, attention, health, privacy, reputation, knowledge, options, independence, and ability to walk away.

Taking care of those resources is not merely protection. It is stewardship.

The same principle eventually applies even to the people we love. We can share what we know. We can make resources available. We can leave doors open. But another person's reading, understanding, agreement, choices, and consequences belong to that person. **Offering does not require chasing. Caring does not require managing.**

And perhaps there is an unexpected return from experiences that once seemed costly.

I may have spent time, attention, generosity, worry, and resources that I would allocate differently today. Those expenditures cannot be recovered. But what they taught me remains mine.

Judgment compounds.

If an experience teaches me to recognize exposure earlier, preserve my resources more carefully, ask better questions, and avoid repeating an expensive mistake, then I emerge with something more valuable than what was lost.

In that sense, I became the ultimate benefactor of my own experience.

The final lesson may therefore be the simplest:

Ask the question that exposes the structure before spending the resource.

Better questions do not merely produce better answers.

They can save years of life.